

Karthik Adari

Dallas, TX | kvkadari@gmail.com | (469) 556-8866 | GitHub | Portfolio | LinkedIn

SKILLS

Data Engineering: Python, R, SQL, ETL/ELT, Snowflake, Cost Optimization, Data Pipelines, Data Modeling, Data Quality, Monitoring/SLAs, Airflow, Databricks, PySpark, Hadoop (HDFS), Hive

Cloud: GCP(BigQuery, Cloud Run, Cloud Functions, Compute Engine, Cloud Storage, Firestore, Cloud Logging, IAM), AWS (Redshift, S3, AWS Glue, Lambda)

Databases & Tools: MongoDB, MySQL, Docker, Git, CI/CD (GitHub Actions), FastAPI

EXPERIENCE

Founder GCP Data Engineer & AI Developer <i>Career in Fortune (AI-Powered Job Platform)</i>	Jan 2025 – Present Dallas, TX
<ul style="list-style-type: none">Led end-to-end architecture of a GCP-native backend (Cloud Run, Firestore, BigQuery) to power job ingestion, search, and analytics workflows.Engineered 30+ production APIs (FastAPI, Cloud Functions) with secure auth and real-time NoSQL data management for platform features.Orchestrated automated ingestion of 100K+ job postings using Playwright + Compute Engine, sustaining 95%+ pipeline reliability.	
Data Engineer Machine Learning Specialist <i>InnovatEV Global</i>	May 2024 – Jan 2025 Dallas, TX
<ul style="list-style-type: none">Built Airflow ETL pipelines to ingest/transform high-volume datasets; added Python data quality checks and managed MongoDB schema evolution.Reduced data issues by 30%+ and cut manual rework by 10+ hours/week through automated validation and pipeline hardening.Containerized services with Docker and deployed on GKE as REST APIs; automated CI/CD with GitHub Actions and added Cloud Logging, reducing release cycles by 60%.	
Data Analyst <i>Corizo</i>	Jan 2021 – Dec 2021 Remote, India
<ul style="list-style-type: none">Automated recurring analytics workflows with scripts, reducing manual effort by 300+ hours annually and improving team throughput.Delivered Tableau KPI dashboards and improved reporting accuracy by 20% through pipeline optimization and cross-functional documentation alignment.	

PROJECTS

HDFS-Based ETL and Reporting on NYC Taxi Dataset <i>Python, Hadoop, HDFS, Hive</i>	Jan 2025 – May 2025
<ul style="list-style-type: none">Ingested and transformed 6.4M+ NYC Yellow Taxi records by converting Parquet to CSV, cleaning data, and loading into HDFS for scalable querying.Created Hive tables and executed analytical SQL to compare Jan 2024 vs Jan 2025 trip volume, fares, and speeds, surfacing seasonal patterns.	
Credit Default Prediction (AmEx Dataset) <i>Python, SQL, XGBoost, Neural Networks</i>	Jan 2024 – May 2024
<ul style="list-style-type: none">Curated and processed 5.5M+ records (190 features) with optimized SQL feature pipelines; achieved 96% accuracy on a sub-10% default-rate dataset.Optimized performance via GridSearch and leveraged XGBoost feature insights to refine a Neural Network, reducing false positives by 25%.	

EDUCATION

The University of Texas at Dallas <i>M.S. in Business Analytics & Artificial Intelligence</i>	Aug 2023 – May 2025 GPA: 3.70/4.00
Indian Institute of Technology Madras <i>Diploma in Programming & Data Science</i>	Jan 2021 – May 2023 GPA: 8.10/10.00
Malaviya National Institute of Technology Jaipur <i>B.Tech in Metallurgical & Materials Engineering</i>	Aug 2019 – May 2023 GPA: 7.46/10.00

CERTIFICATIONS

Google Cloud Professional Data Engineer
Salesforce Certified Associate